

A numerical audit, with known errors

Synthetic demonstration - not a real scientific paper

This document intentionally contains correct and incorrect claims. It tests source tracking, arithmetic and abstention. The values are invented solely as labelled test data.

Methods

Context demo/test/run1 denotes the Demo dataset, held-out test split, first run, with the same classification setting for both models. The table below is machine-readable and has explicit context labels.

Results

Model	Accuracy (%)	Context
Model A	87.4	demo/test/run1
Model B	84.2	demo/test/run1

[context: demo/test/run1] Model A accuracy = 87.4%.

[context: demo/test/run1] Model B accuracy = 84.2%.

A separate arithmetic example uses a baseline accuracy of 80.0% and a new accuracy of 92.0%. These values are not the model-comparison experiment above.

Discussion

[context: demo/test/run1] Model B achieved higher classification accuracy than Model A.

[context: demo/test/run1] Model A achieved higher accuracy than Model B.

Accuracy improved by 12% from 80.0% to 92.0%.

Accuracy improved by 15% from 80.0% to 92.0%.

Accuracy increased by 12 percentage points from 80.0% to 92.0%.

Our findings prove that the approach generalizes to all hospitals.

Limitations

Only one synthetic dataset and one run are described. No independent replication or external validity evidence is supplied. The broad hospital claim has no supporting experiment.

Expected audit behavior

Flag the claim that Model B has higher accuracy. Distinguish the incorrect 12% relative change from the correct 15% relative change and 12 percentage-point change. Abstain on the broad generalization. Preserve the actual PDF page and evidence identifiers for each finding.